

A Appendix

S1.1 Detailed Knowledge Graph Generation

In this section, we explain each step in our graph generation process, illustrating how unstructured chemical text is transformed into a structured representation suitable for downstream tasks.

S1.1.1 Text Preprocessing

Our preprocessing pipeline is designed to isolate the most informative portions of ChemRxiv articles and prepare them for entity and relation extraction. First, we retrieved all ChemRxiv articles whose licenses permit redistribution via the ChemRxiv API. Using regular expressions, we extracted the introduction section of each paper, as this typically focuses on concise, objective statements about chemical phenomena. From each introduction, we extracted the first 500 words—this limit balances coverage of key background information against the risk of including less relevant or overly general text. To avoid fragmenting paragraphs, we ensured that no paragraph was split across chunks. Finally, we segmented each introduction into contiguous chunks of up to 128 words, preserving natural sentence boundaries. This segmentation reduces the input length for subsequent processing steps, ensuring that each chunk remains within the token limits of our downstream models while retaining semantic coherence.

S1.1.2 Node Extraction

Once the text is segmented, we identify chemical entities and their interrelations. We apply a named entity recognition (NER) model based on the PubMedBERT architecture, fine-tuned on multiple chemical datasets, to detect candidate entities within each 128-word chunk. To further improve precision and eliminate spurious mentions, we refine the NER output using OpenAI’s gpt-4o, prompting it with a few shots to verify whether each detected span indeed corresponds to a chemically meaningful entity and to standardize entity labels (for instance, converting “MeOH” to “methanol”). Below, we showed the prompt for Entity verification.

You are a chemistry expert specializing in entity recognition. Your task is to **validate**

and filter the extracted entities, ensuring they are **chemically meaningful** based on the provided text. Remove any irrelevant terms, including general descriptors, numerical values, reaction conditions, and vague terms.

Entities Extracted by NER:

{entities}

Text for Context:

{text}

Criteria for Valid Entities:

- 3 Chemical compounds (e.g., "HCl", "Sodium hydroxide", "Ethanol", "Benzene")
- 3 Chemical elements (e.g., "Carbon", "Oxygen", "Cesium")
- 3 Specific catalysts, solvents, reagents (e.g., "Cs₂CO₃", "Toluene", "Palladium")

Remove the Following Types of Entities:

- × Generic terms (e.g., "Reaction", "Solvent", "Acid", "Base", "Solution")
- × Experimental conditions (e.g., "pH", "Temperature", "2 M", "Strong acid")
- × Measurement terms (e.g., "X-ray diffraction", "NMR")
- × General descriptors (e.g., "High concentration", "Low efficiency")

Output Format:

Return only a **Python list** of valid chemical entities, with no explanations, markdown, or extra formatting.

446 S1.1.3 Edge Extraction

447 After obtaining a vetted set of entities, we extract pairwise relations using the same gpt-4o
 448 instance. For each pair of entities co-occurring within a chunk, we prompt the model with a
 449 few shots to classify or generate the nature of their relationship, producing triplets of the
 450 form (entity_A, relation, entity_B). The result of this step is a set of entity nodes (chemical
 451 names) and directed edges (relation labels) extracted directly from the text, forming a raw
 452 triplet collection that reflects the functional associations present in the chemical literature.
 453 Below we showed the prompt for Edge Extraction.

You are an expert in chemical text analysis. Your task is to extract **only chemically**

meaningful relationships between a given set of entities from the provided text.

Guidelines for Relation Extraction:

1. **Entity Matching:** Consider only the entities provided in the given set. If an entity appears in the text but has no meaningful chemical relationship with another entity in the set, ignore it.
2. **Chemically Significant Relations Only:** Extract relations that describe actual **chemical interactions, transformations, or properties** (e.g., "reacts with," "catalyzes," "dissolves in," "produces").
3. **Factual Relations:** Only extract factual relations. Avoid observations, opinions, and findings.
4. **Tuple Format:** Output extracted facts in the form of **(entity1, relation, entity2)**.
5. **Avoid Generic Relations:** Exclude weak relations like "is," "are," "exists," "relates to." Focus on **specific interactions**.

Valid Relation Types (Examples):

- 3 "reacts with"
- 3 "catalyzes"
- 3 "binds to"
- 3 "dissolves in"
- 3 "oxidizes"
- 3 "inhibits"
- 3 "precipitates with"
- 3 "acts as a solvent for"
- 3 "is synthesized from"

Avoid These Weak Relations: Exclude relations such as "is," "are," "has," "exists."

Entities Provided:

{entities}

Text:

{text}

Extract at most {max_facts} factual statements.

Output Format:

Provide the output as a **Python list of tuples**, containing only the extracted relationships without any code formatting, backticks, or markdown.

Example Output:

– [("HCl", "dissolves in", "Water"), ("HCl", "reacts with", "Sodium hydroxide")]

454 **S1.1.4 Knowledge Enrichment**

455 To enrich each node with descriptive metadata and resolve naming ambiguities, we retrieved
 456 supplemental information from two external resources: Wikipedia and PubChem. From
 457 Wikipedia, we extracted the introductory summary of each entity’s page, providing a
 458 concise description of its common usage, historical context, or primary function. From
 459 PubChem, we obtained the official compound name (Record Title) alongside additional
 460 names and identifiers drawn from the “Names and Identifiers” section. We retrieved a
 461 textual description from the “Record Description” heading, summarizing key properties
 462 or applications. Safety annotations, including hazard statements or pictograms, were
 463 collected from the “Chemical Safety” subsection. Structural information was captured in
 464 the form of the canonical SMILES string under “Computed Descriptors,” and the molecular
 465 formula was fetched from the “Molecular Formula” field. Finally, we extracted computed
 466 physicochemical properties—such as molecular weight, topological polar surface area,
 467 and log P values—from the “Computed Properties” list within “Chemical and Physical
 468 Properties.” All of these metadata fields were stored as additional text and added as the
 469 external info for each of the nodes.

470 **S1.1.5 Graph Generation**

471 At this stage, we constructed the graph by forming triplets of the form (node, edge, node).
 472 Each node was linked to the original text segment from which it was extracted and, if
 473 available, to its corresponding external metadata. Likewise, each edge was associated with
 474 the specific text chunk that produced it. In this way, both nodes and edges maintain direct
 475 references to their source text, ensuring traceability throughout the knowledge graph.

476 **S1.2 Detailed Question Generation**

477 **S1.2.1 Path Sampling from the Knowledge Graph**

478 To generate multi-hop questions, we first sampled paths of varying lengths from the
 479 constructed knowledge graph using a randomized breadth-first search (BFS) path sampling
 480 algorithm. During path sampling, we enforced that each edge in a sampled path originates
 481 from a distinct source text, thereby encouraging the integration of information from multiple,
 482 separate context passages. Concretely, a path of length K comprises $K + 1$ entities and
 483 traces through K different source texts extracted from the original ChemRxiv database. By
 484 requiring unique sources for each edge, we ensure that correct solutions must draw on
 485 evidence scattered across several documents rather than a single paragraph.

486 **S1.2.2 One-Hop Question Formulation**

487 Adopting a bottom-up approach, we generated individual one-hop questions for each edge
 488 along a sampled path. For every triplet (entity₁, relation, entity₂), we framed a question
 489 whose answer is entity₁ by asking, “Which entity holds the relation relation to entity₂?” If
 490 the initial phrasing lacked sufficient specificity or clarity, we invoked a language model
 491 (OpenAI’s o3-mini) to augment the prompt with metadata drawn from the original text
 492 segment, such as contextual phrases or qualifying details. This enrichment step ensures that
 493 each one-hop question remains clear, precise, and answerable from the associated source
 494 text alone. Below is the prompt for generating the one hop questions.

You are given a text along with an entity and its relation to another entity.

Entity 1: {entity1}
 Relation: {relation}
 Entity 2: {entity2}
 Text: {text}

Information about Entity1: {entity1_meta if entity1_meta else None}

Your task is to generate a factual question whose answer is Entity1.
 The question should ask for the entity that has the specified relation to Entity2.
 Do not mention the answer (which is Entity1) in the question.
 Ensure that the question is factual and can be answered solely based on the given text and the information about Entity1.
 Do not refer to sections such as "Abstract," "Table 1," "in the text," or "in the article."

If Entity1 and relation are not specific enough (i.e., multiple answers are possible), add descriptions from the text or from the information about Entity1 to make it specific so that Entity1 is the only answer.

Return a dictionary without any code formatting, backticks, or markdown, with keys \q" and \a".

495 S1.2.3 Multi-Hop Question Aggregation

496 Once all one-hop questions for a path were vetted, we combined them into a single multi-hop
 497 question by prompting OpenAI's o3-mini model, with a few shots. The final aggregated
 498 question is constructed by starting with the sub-question corresponding to the last edge
 499 (i.e., the entity nearest to the "tail" of the path) and then chaining backward through each
 500 preceding entity until reaching entity₁ of the first relation. In other words, if the sampled
 501 path is

$$(\text{entity}_1 \xrightarrow{\text{relation}_1} \text{entity}_2), (\text{entity}_2 \xrightarrow{\text{relation}_2} \text{entity}_3), \dots, (\text{entity}_K \xrightarrow{\text{relation}_K} \text{entity}_{K+1}),$$

502 the aggregated question asks first about entity_{K+1} (the final tail), then uses its answer as
 503 context for the penultimate question, and so on, so that the final answer corresponds to
 504 entity₁. This reverse-chaining structure ensures that the multi-hop prompt leads directly to
 505 the original target node while preserving logical flow. Below we showed the prompt for
 506 generating the multi hop question.

You are given multiple factual questions and their answers that are logically connected.

Your task is to chain them into a single, coherent multi-hop question that requires multiple reasoning steps.
 Ensure that the (only) answer is the answer to the first question, and the question naturally follows from the facts given.
 You have to start from the last generated question and build up a single multi-hop question so it aggregates them all and the answer is the answer to the first question.
 None of the answers to any of the questions should be in the generated question.

Here is an example:

Example:

Q1: What is oxidized to form Carbon Dioxide?

A1: Methane

Q2: What is used in Photosynthesis?

A2: Carbon Dioxide

Q3: What produces Oxygen?

A3: Photosynthesis

Multi-hop question:

Q: What is oxidized to produce a substance that is used in a process that results in Oxygen?

A: Methane

Here are the generated questions and answers:

{formatted_qas}

Return a python dictionary without any code formatting, backticks, or markdown, with keys "q" (multi-hop question) and "a" (final answer).

507 **S1.2.4 Verification and Filtering**

508 During verification, we first reviewed each one-hop question for clarity, chemical relevance,
 509 and direct answerability based on its corresponding source text. Below the prompt for
 510 one-hop question evaluation is shown.

You are a chemistry expert. Your task is to determine if the given question is a factual

chemistry question, unambiguous (has only one answer), and answerable based on the provided context. A factual question must be based on actual chemical properties, reactions, or experimentally verified principles and must be strictly related to chemistry. An answerable question should be solvable based on the given context and must not be open-ended or have multiple correct answers. MAKE SURE THE QUESTION HAS ONLY ONE CORRECT ANSWER. There shouldn't be any other entity except for the given answer that could be another answer.

Question:
{question}

Answer:
{answer}

Context:
{context}

Please analyze the context and verify if the question is factual, unambiguous, and answerable. If the question is factual, has only one correct answer, is strictly related to chemistry, and can be answered based on the context, return "yes." Otherwise, return "no."

Examples of Factual Chemistry Questions:

"What dissolves in water and evaporates at 0 °C?"

"What catalyst is used in the reaction between A and B?"

Examples of Non-Factual or Ambiguous Chemistry Questions:

"What is the song of Nirvana that is a chemical entity?"

"What chemical entity and structural unit form the layered hydroxide structures with intercalated water ions used in battery materials and OER catalysis?" (M(OH)₆ and -Ni(OH)₂ are valid answers)

Questions that have multiple possible correct answers or are not strictly related to chemistry.

511 Next, the multi-hop question underwent an additional evaluation step to confirm that
512 the logical chain—from the final sub-question back to the first—correctly guides a reader
513 (or model) to entity₁. We employed an LLM-based verification process to assess factual
514 accuracy, answerability given available context and metadata, and overall coherence
515 among sub-questions. Feedback from domain experts was continuously incorporated
516 into the prompt templates to refine verification accuracy. Any questions—either one-hop or
517 multi-hop—that were answered incorrectly by all evaluated models were excluded from
518 the benchmark to minimize ambiguity and ensure high-quality, unambiguous reasoning
519 tasks. Below the prompt for evaluating the whole path is shown

You are a chemistry expert. Your task is to determine if the given question is a factual

chemistry question and answerable based on the provided path.

Path Information:
{path_text}

Question:
{question}

Answer:
{answer}

Please analyze the path and verify if the question is a factual chemistry question and can be answered based on the given path. A factual question must be based on actual chemical properties, reactions, or experimentally verified principles. An answerable question should be solvable based on the given path. If the question is factual and answerable, return "yes". If it contains speculation, opinions, or lacks verifiable chemical grounding, or it is not solvable, return "no".

Examples of Factual Chemistry Questions:

"What dissolves in water?"

"What catalyst is used in the reaction between A and B?"

"Which compound undergoes oxidation in this reaction?"

"What product is formed when sodium reacts with chlorine?"

Examples of Non-Factual Chemistry Questions:

"Why do some scientists think this reaction is inefficient?"

"What is the best solvent for this reaction?"

"Is this reaction useful in industry?"

"Do you think this compound is a good catalyst?"

Provide only "yes" or "no" as your response.

520 **S1.2.5 Short-Answer Design and Rationale**

521 To minimize the impact of writing style and summarization on accuracy evaluation, all
522 questions were deliberately designed to elicit short, concise answers. Answering a multi-hop
523 question requires decomposing it into its constitutive one-hop sub-questions, retrieving
524 each intermediate answer, and then combining them to arrive at the final answer. Even
525 when full context is available, a correct response cannot be produced if a model fails to
526 infer and integrate multiple pieces of knowledge. By keeping answers brief and focusing
527 on discrete factual steps, we ensure that performance evaluation reflects a model's ability
528 to conduct multi-hop inference rather than its capacity to paraphrase or generate lengthy
529 explanations.

530 **S1.3 Rejected Questions**

531 There were a couple of reasons that we identified for rejecting of some questions from
532 the benchmark. One was having multiple valid answers, and second was the answer was
533 present in the question. Some examples

Q1:

Context:

Researchers have developed an anode material based on NiCo_rGO (Nickel–Cobalt–reduced Graphene Oxide). In some variants, the NiCo_rGO is further decorated with palladium (Pd) nanoparticles to enhance catalytic performance.

Question:

Which component in the electrode structure functions as a catalyst at the anode when incorporating decorated NiCo_rGO?

Issue:

The question is declined due to ambiguity — two distinct answers are technically correct based on the variant of the material:

- If the material is **NiCo_rGO** without Pd, the catalyst is **Nickel (Ni)**.
- If the material is **Pd-decorated NiCo_rGO**, the catalyst is **Palladium (Pd)**.

Since the phrasing of the question does not clearly disambiguate which material variant is being used, it leads to multiple valid interpretations. Therefore, it cannot be accepted as a single-answer question.

Q2:**Context:**

Hypervalent iodine compounds such as diaryliodonium salts are widely used as electrophilic arylation reagents. According to the source text, these salts are employed in both **transition metal-catalyzed** and **metal-free** arylation reactions. These reactions can be used to functionalize aromatic compounds, including halogen-substituted analogues like CIMPPC, by replacing hydrogen atoms.

Question:

Which type of arylation reaction that employs electrophilic arylation reagents utilizes diaryliodonium salts in hypervalent iodine chemistry?

Expected Answer:

Transition metal-catalyzed arylations

Reason for Rejection:

The question is declined due to the presence of **multiple valid answers**. The source explicitly states that diaryliodonium salts are used in:

- **Transition metal-catalyzed arylations**, and
- **Metal-free arylations**.

Both are equally valid interpretations of the question. Without further constraints or clarification, the question has more than one correct answer and does not meet the single-answer requirement.

534 S1.4 Detailed Performance Based on Context Availability

535 Figure S1 shows that model performance is strongly influenced by whether context is
 536 provided in the input. In particular, Claude 3.7 with extended thinking achieved a
 537 correctness rate of 84% with context, whereas o3-mini recorded the highest correctness
 538 rate (48%) when the context was absent. Note that o3-mini was primarily used to generate
 539 the questions, which may have introduced a slight bias resulting in its minor improvement
 540 in correctness. Figures S2 and S3 illustrate the token usage and latency of the models,
 541 respectively.